

Next Generation Datacenter Operating Systems

A general-purpose
OS/application interface
for programmable hardware

Datacenter applications have demanding requirements.

- High perf (low latency, high throughput)
- Low tail latency
- Virtualization/multi-tenant sharing
- Security isolation
- Performance isolation
- Flexibility

To help meet those requirements, datacenter servers have added lots of (programmable) hardware.

- IO MMU
- Virtualized/Programmable NIC
- Programmable (remote) SSD
- (Programmable) RDMA

Which has caused the systems stack to change at a rapid pace.

- The hardware keeps changing (e.g., new NICs, new features).
- The hardware/software interface keeps changing (e.g., let's put everything into the NIC! Let's not!).
- New systems interfaces keep appearing.

How to build a demanding datacenter app in this world?

- Modify applications for every new technology/system.
- Require new systems to support legacy interfaces (e.g., POSIX)
- Design new general systems/application interface for programmable hardware

Requirements for new systems interface

- Separate app from hardware to allow hardware to evolve independently
- Efficient for transferring data (not a serializing interface!)
- Efficient to implement in hardware or software

Everything is a file

zero-copy queue

- Replaces file descriptors and pipes with queue descriptors and queues
- Is zero-copy for efficient app processing without cache pollution
- Offers insight into granularity for filtering, merging, sorting

Interface

```
qid = socket(domain, type, protocol)
```

```
qid = open(file)
```

```
qid = accept(qid)
```

```
insert(qid, *sga)
```

```
*sga = dequeue(qid)
```

```
qid = filter(qid, *filter_func)
```

```
qid = merge(qid, qid)
```

```
qid = sort(qid, *sort_func)
```

```
(sga = scatter gather array = list of bufs)
```

Use cases

- File/storage server
- Memcached
- Meta-data server
- Replicated service
- Others?

Available Hardware (MSR only)

	programmable?	programming mode	features
RDMA	static		direct memory access
RDMA	partially programmable	firmware changes	
RDMA	fully programmable	FPGA	
SSD	fully programmable	start-up	SR-IOV (end 2018)
SSD	remote access	FPGA	

Fast Context Switching for Fine-grained Process Scheduling

High-performance datacenter apps make poor use of CPUs

- Existing solution is to pin threads, which under-utilizes the CPU for bursty workloads.
- No multi-tenancy to even out bursts.
- CPU to go into a lower power mode in between bursts, increasing tail latency

New hardware has lowered the cost of context switches

- Faster switches between rings
- Tagged TLBs
- Partitioned caches

Making it feasible to schedule the CPU for shorter periods.

Datacenter applications have natural interrupt points

- High-performance datacenter apps are typically request processors
- Less data shared between requests (measure this!)
- OS has no insight into these points

Cooperative scheduling will perform
better than interrupts or polling

Request-based process scheduling

- Can potentially be done without modifying app (e.g., changing libevent)
- Allow even high performance apps to share a CPU
- Lower (tail) latency for everyone
- Better performance isolation

Exactly once RPC hardware